

# Class 6

Linh Tran (A16435846)

Functions are how we get work done in R. We call functions to do everything from reading data to doing analysis and outputting plots and results.

All functions in R have at least 3 things:

- a **name** (you get to pick this)
- input **arguments** (there can be only one or loads - again your call)
- the **body** (where the work gets done, this code between the curly brackets)

## A first silly function

Let's write a function to add some numbers. We can call it `add()`

```
x <- 10  
y <- 10  
x+y
```

```
[1] 20
```

```
add <- function(x) {  
  y <- 10  
  x+y  
}
```

Can I just use my new function?

```
add(1)
```

```
[1] 11
```

Let's make it

```
add <- function(x, y=1) {  
  x+y  
}  
add(x=10,y=10)
```

```
[1] 20
```

```
add (10)
```

```
[1] 11
```

```
add(10,100)
```

```
[1] 110
```

## 2nd example grade() function

Write a function to grade student work

We will start with a simple version of the problem and the following example student vectors:

```
# Example input vectors to start with  
student1 <- c(100, 100, 100, 100, 100, 100, 100, 90)  
student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)  
student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)
```

Start with student 1

```
mean(student1)
```

```
[1] 98.75
```

Use na.rm, rm stands for *remove*

```
mean(student2, na.rm=TRUE)
```

```
[1] 91
```

```
mean(student3, na.rm=TRUE)
```

```
[1] 90
```

Ok let's try to work with student1 and find (and drop) the lowest score.

```
student1
```

```
[1] 100 100 100 100 100 100 100 90
```

Google told me about min() and max()

```
min(student1)
```

```
[1] 90
```

```
which.min(student1)
```

```
[1] 8
```

```
student1[1:4]
```

```
[1] 100 100 100 100
```

```
student1[8]
```

```
[1] 90
```

```
student1[which.min(student1)]
```

```
[1] 90
```

```
student1[-8]
```

```
[1] 100 100 100 100 100 100 100
```

Our first working snippet that drops the lowest score and calculates the mean

```
mean(student1[-which.min(student1)])
```

```
[1] 100
```

```
x <- student3  
mean(x[-which.min(x)], na.rm=T)
```

```
[1] NaN
```

Our approach to the NA problem (missing homeworks): We can replace all NA values with zero.

1st task is find the NA values (i.e. where are they in the vector)

```
x==90
```

```
[1] TRUE  NA  NA  NA  NA  NA  NA  NA
```

```
x <- student2  
mean(x[-which.min(x)], na.rm=T)
```

```
[1] 92.83333
```

```
is.na(x)
```

```
[1] FALSE  TRUE FALSE FALSE FALSE FALSE FALSE
```

```
x[is.na(x)] <- 0
```

I have found the NA (TRUE) values from `is.na()` now I want to make them equal to zero (overwrite them/mask them etc.)

For example,

```
y <- 1:5  
y
```

```
[1] 1 2 3 4 5
```

```
y[y>3] <- 0  
y
```

```
[1] 1 2 3 0 0
```

To determine the overall grade and drop the lowest score. I want to combine the `is.na(x)` with making these elements equal to zero. And then take this “masked” (vector of student scores with NA values as zero) and drop the lowest to get the mean

```
x <- student2  
x[is.na(x)] <- 0  
mean(x[-which.min(x)])
```

```
[1] 91
```

Now I can turn my most awesome snippet into my first function

```
grade <- function(x){  
  x[is.na(x)] <- 0  
  mean(x[-which.min(x)])  
}
```

```
grade(student1)
```

```
[1] 100
```

```
grade(student2)
```

```
[1] 91
```

```
grade(student3)
```

```
[1] 12.85714
```

## R Functions Lab (Class 06):

Q1. Write a function `grade()` to determine an overall grade from a vector of student homework assignment scores dropping the lowest single score. If a student misses a homework (i.e. has an NA value) this can be used as a score to be potentially dropped.

```
url <- "https://tinyurl.com/gradeinput"
gradebook <- read.csv(url, row.names=1)
head(gradebook)
```

	hw1	hw2	hw3	hw4	hw5
student-1	100	73	100	88	79
student-2	85	64	78	89	78
student-3	83	69	77	100	77
student-4	88	NA	73	100	76
student-5	88	100	75	86	79
student-6	89	78	100	89	77

The ‘`apply()`’ function in R is super useful but can be a little confusing to begin with. Let’s have a look how it works

```
ans <- apply(gradebook, 1, grade)
ans
```

student-1	student-2	student-3	student-4	student-5	student-6	student-7
91.75	82.50	84.25	84.25	88.25	89.00	94.00
student-8	student-9	student-10	student-11	student-12	student-13	student-14
93.75	87.75	79.00	86.00	91.75	92.25	87.75
student-15	student-16	student-17	student-18	student-19	student-20	
78.75	89.50	88.00	94.50	82.75	82.75	

Q2. Using your `grade()` function and the supplied gradebook, Who is the top scoring student overall in the gradebook?

```
which.max(ans)
```

```
student-18
18
```

```
max(ans)
```

```
[1] 94.5
```

Q3. From your analysis of the gradebook, which homework was toughest on students (i.e. obtained the lowest scores overall?

```
toughest <- which.min(apply(gradebook,2,mean, na.rm= TRUE))
toughest
```

```
hw3
```

```
3
```

Q4. Optional Extension: From your analysis of the gradebook, which homework was most predictive of overall score? want the students that score high to score high

```
#ans
cor(gradebook$hw1, ans)
```

```
[1] 0.4250204
```

```
#ans
cor(gradebook$hw5, ans)
```

```
[1] NA
```

```
gradebook$hw5
```

```
[1] 79 78 77 76 79 77 100 100 77 76 100 100 80 76 NA 77 78 100 79
[20] 76
```

Make all NA values into zero

```
mask <- gradebook
mask[is.na(mask)] <- 0
```

```
cor(mask$hw5,ans)
```

```
[1] 0.6325982
```

Now we can use ‘`apply()`’ to examine the correlation of every assignment in the masked grade-booJ to the overall score for each student in the class.

```
apply(mask,2,cor, y=ans)
```

	hw1	hw2	hw3	hw4	hw5
	0.4250204	0.1767780	0.3042561	0.3810884	0.6325982

Hw5 was the most predictive of overall score.